



Introduction to Deep Learning

Recurrent Neural Networks for Sequential Data (Time Series)

MATH 370: Machine Learning


Tanujit Chakraborty

@ Sorbonne


ctanujit@gmail.com

Learning from Time-Series Data

- The input is a sequence of (non-i.i.d.) examples y_1, y_2, \dots, y_t .
- The problem may be supervised or unsupervised, e.g.,
 - Forecasting: Predict y_{t+1} using y_1, y_2, \dots, y_t
 - Cluster the examples or perform dimensionality reduction / Anomaly detection
- Evolution of time-series data can be attributed to several factors




- Teasing apart these factors of variation is also an important problem.



Auto-regressive Models

- **Auto-regressive (AR):** Regress each example on p previous lagged values - AR(p) model

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$



Auto-regressive Model (shown above: 2nd order AR)

- **Moving Average (MA):** Regress each example on q previous stochastic errors - MA(q) model
- **Auto-regressive Integrated Moving Average (ARMA):** Regress each example of p previous lagged values and q previous stochastic errors


$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

where y_t' is the differenced series (if the data is nonstationary and differencing is applied). We call this an **ARIMA**(p, d, q) model.

$$y_t' = c + \phi_1 y_{t-1}' + \cdots + \phi_p y_{t-p}' + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t,$$

State-Space Models

- Assume that each observation y_t in the time-series is generated by a low-dimensional latent factor x_t (one-hot or continuous)



State-Space Model (shown above: 1st order SSM)

- Basically, a generative latent factor model: $y_t = g(x_t)$ and $x_t = f(x_{t-1})$, where g and f are probability distributions.
- Some popular **SSMs**: Hidden Markov Models (one-hot latent factor x_t), Kalman Filters (real-valued latent factor x_t)
- **Note:** Models like RNN/LSTM are also similar, except that these are not generative (but can be made generative)



Long Memory and ARFIMA Process

- **Definition (Long Memory).** Let $\{X_t, t \in \mathbb{Z}\}$ be a weakly stationary univariate process with auto-covariance function $\gamma_X(k)$ and spectral density function

$$f_X(\lambda) = (2\pi)^{-1} \sum_{k=-\infty}^{\infty} \gamma_X(k) \exp(-ik\lambda) \text{ for } k \in \mathbb{Z} \text{ and } \lambda \in [-\pi, \pi].$$

Then, $\{X_t\}$ has long memory if $\sum_{k=-\infty}^{\infty} |\gamma_X(k)| = \infty$, and short memory otherwise.

Equivalently, as $|\lambda| \rightarrow 0$, $f_X(\lambda) \rightarrow \infty$ for long memory,

- The most popular long-memory model for level data y_t is the ARFIMA(p, d, q) model introduced by Granger and Joyeux (1980) and Hosking (1981). Specifically, an ARFIMA(p, d, q) process y_t is defined by

$$(1 - B)^d y_t = x_t$$

- B is the one-dimensional backshift operator, x_t is an ARMA(p, q) process that captures short-range dependence, and d is a fractional differencing parameter.
- Typically, d is chosen such that $-1/2 < d < 1/2$ to ensure that y_t is stationary and invertible.

Fractional Differencing



Biometrika (1981), **68**, 1, pp. 165–76
Printed in Great Britain

165

Fractional differencing

BY J. R. M. HOSKING

Institute of Hydrology, Wallingford, Oxfordshire

SUMMARY

The family of autoregressive integrated moving-average processes, widely used in time series analysis, is generalized by permitting the degree of differencing to take fractional values. The fractional differencing operator is defined as an infinite binomial series expansion in powers of the backward-shift operator. Fractionally differenced processes exhibit long-term persistence and antipersistence; the dependence between observations a long time span apart decays much more slowly with time span than is the case with the more commonly studied time series models. Long-term persistent processes have applications in economics and hydrology; compared to existing models of long-term persistence, the family of models introduced here offers much greater flexibility in the simultaneous modelling of the short-term and long-term behaviour of a time series..





Figure: Fractional Differencing applied to DAX index

Recurrent Connections in Deep Neural Networks




- Feedforward nets such as MLP assume independent observations



Feedforward neural networks are not ideal when inputs $[x_1, x_2, \dots, x_N]$ and/or outputs $[y_1, y_2, \dots, y_N]$ represent sequential data (e.g., sequence of words, video (sequence of frames), etc.)




- A **recurrent structure** can be helpful if each input and/or output is a sequence



RNNs

- RNNs are used when each input or output or both are **sequences of tokens**




If the input is a word sequence, then each x_n represent the corresponding word's embedding (either a pre-computed word embedding like word2vec or a learned word embedding)

- Hidden state h_t is supposed to remember everything up to time $t - 1$. However, in practice, RNNs have difficulties remembering the distant past
 - Variants such as LSTM, GRU, etc mitigate this issue to some extent
- Slow processing is another major issue (e.g., can't compute h_t before computing h_{t-1})

Recurrent Neural Networks



- A basic RNN's architecture (assuming input and output sequence have same lengths)




- RNN has three sets of weights $\mathbf{W}, \mathbf{U}, \mathbf{v}$
- \mathbf{W} and \mathbf{U} model how h_t at step t is computed: $\mathbf{h}_t = g(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1})$
- \mathbf{v} models the hidden layer to output mapping, e.g., $\mathbf{y}_t = o(\mathbf{v}\mathbf{h}_t)$
- **Important:** Same $\mathbf{W}, \mathbf{U}, \mathbf{v}$ are used at all steps of the sequence (weight sharing)

g is some activation function like **ReLU**


o depends on the nature of y_t . If it is categorical then o can be **softmax**

Recurrent Neural Nets (RNN)

- A more “micro” view of RNN (the transition matrix U connects the hidden states across observations, propagating information along the sequence)



RNN in Action




Workflow of RNN:

- The gif above reflects the magic of recurrent networks.
- It depicts 4 timesteps. The first is exclusively influenced by the input data.
- The second one is a mixture of the first and second inputs. This continues on.
- You should recognize that, in some way, network 4 is "full".
- Presumably, timestep 5 would have to choose which memories to keep and which ones to overwrite.
- This is very real. It's the notion of memory "capacity".
- As you might expect, bigger layers can hold more memories for a longer period of time.

Training RNN



- Trained using **Backpropagation Through Time** (forward propagate from step 1 to end, and then backward propagate from end to step)
- Think of the time-dimension as another hidden layer and then it is just like standard backpropagation for feedforward neural nets




- They learn by fully propagating forward from 1 to 4 (through an entire sequence of arbitrary length), and then backpropagating all the derivatives from 4 back to 1.
- You can also pretend that it's just a funny shaped normal neural network, except that we're re-using the same weights (synapses 0,1, and h) in their respective places.
- Other than that, it's normal backpropagation.

Pic source: <https://iamtrask.github.io/>

RNN Applications



- In many problems, each input, each output, or both may be in form of sequences




- Different inputs or outputs need not have the same length
- Some examples of prediction tasks in such problems
 - Image captioning:** Input is image (not a sequence), output is the caption (word sequence)
 - Document classification:** Input is a word sequence, output is a categorical label
 - Machine translation:** Input is a word sequence, output is a word sequence (in different language)
 - Stock price prediction:** Input is a sequence of stock prices, output is its predicted price tomorrow
 - No input – just output (e.g., **generation** of random but plausible-looking text)


Recurrent Neural Networks: Some Examples



- Consider generating a sequence y_1, y_2, \dots, y_T given an input x




- Predicting the sentiment of a movie review




Recurrent Neural Networks: Some Examples



- Parts of speech tagging (or “aligned” translation; input and output have same length)




- “Unaligned” translation (input and output can have different lengths)



- In the unaligned case, generation stops when an “end” token (e.g., <END>) is generated on the output side

Recurrent Neural Networks: Some Examples

- Unconditional generation (no input, only an output sequence is generated given a RNN that was trained using some training data containing several sequences)




- Each generated word/token is fed to the next step’s hidden state
- Generation stops when an “end” token (e.g., <END>) is generated

For RNNs, Long Distant Past is Hard to Remember



- The hidden layer nodes h_t are supposed to summarize the past up to time $t - 1$




- In theory, they should. In practice, they can't. Some reasons
 - Vanishing gradients along the sequence too (due to repeated multiplications)
 - past knowledge gets “diluted”
 - Hidden nodes also have limited capacity because of their finite dimensionality
- Various extensions of RNNs have been proposed to address forgetting
 - Gated Recurrent Units (GRU), Long Short Term Memory (LSTM) (Hochreiter and Schmidhuber, mid-90s)

GRU and LSTM

- Essentially an RNN, except that the hidden states are computed differently
- Recall that RNN computes the hidden states as

$$\mathbf{h}_t = \tanh(\mathbf{W}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1})$$

- For RNN: State update is multiplicative (weak memory and gradient issues)
- GRU and LSTM contain specialized units and “memory” which modulate what/how much information from the past to retain/forget




σ FC layer with Activation function

\times Element-wise Operator

↑ Copy

↗ Concatenate

Pic source: <https://d2l.ai/>



σ FC layer with Activation function


\times Element-wise Operator

↑ Copy

↗ Concatenate

Capturing Long-Range Dependencies

- Idea: Augment the hidden states with **gates** (with parameters to be learned)
- These gates can help us remember and forget information “**selectively**”



Pic source: www.deeplearning4j.org

- The hidden states have 3 type of gates: Input (bottom), Forget (left), Output (top)
- Open gate denoted by 'o' closed gate denoted by '-'



LSTM

- In contrast, LSTM maintains a “context” C_t and computes hidden states as

$$\begin{aligned}\hat{C}_t &= \tanh(\mathbf{W}^c \mathbf{x}_t + \mathbf{U}^c \mathbf{h}_{t-1}) && (\text{“local” context, only up to immediately preceding state}) \\ i_t &= \sigma(\mathbf{W}^i \mathbf{x}_t + \mathbf{U}^i \mathbf{h}_{t-1}) && (\text{how much to take in the local context}) \\ f_t &= \sigma(\mathbf{W}^f \mathbf{x}_t + \mathbf{U}^f \mathbf{h}_{t-1}) && (\text{how much to forget the previous context}) \\ o_t &= \sigma(\mathbf{W}^o \mathbf{x}_t + \mathbf{U}^o \mathbf{h}_{t-1}) && (\text{how much to output}) \\ C_t &= C_{t-1} \odot f_t + \hat{C}_t \odot i_t && (\text{a modulated additive update for context}) \\ h_t &= \tanh(C_t) \odot o_t && (\text{transform context into state and selectively output})\end{aligned}$$

- Note: \odot represents elementwise vector product. Also, state updates now additive, not multiplicative. Training using backpropagation through time.
- Many variants of LSTM exists, e.g., using C_t in local computations, Gated Recurrent Units (GRU), etc. Mostly minor variations of basic LSTM above.

Do LSTM really have long memory? (ICML'2020)



Do RNN and LSTM have Long Memory?

Jingyu Zhao¹ Feiqing Huang¹ Jia Lv² Yanjie Duan² Zhen Qin² Guodong Li¹ Guangjian Tian²

Abstract

The LSTM network was proposed to overcome the difficulty in learning long-term dependence, and has made significant advancements in applications. With its success and drawbacks in mind, this paper raises the question - do RNN and LSTM have long memory? We answer it partially by proving that RNN and LSTM do not have long memory from a statistical perspective. A new definition for long memory networks is further introduced, and it requires the model weights to decay at a polynomial rate. To verify our theory, we convert RNN and LSTM into long memory networks by making a minimal modification, and their superiority is illustrated in modeling long-term dependence of various datasets.




Figure: Autocorrelation plot of traffic and DJI datasets
(To visualize the long memory in the dataset)

Memory RNN and Bidirectional RNN

- RNNs and GRU and LSTM only remember the information from the previous tokens
- Memory RNN and Bidirectional RNN** can remember information from the past and future tokens





Figure 1. The MRNN cell structure.


Ref: <https://proceedings.mlr.press/v119/zhao20c>



Ref: <https://ieeexplore.ieee.org/document/650093>

Exercise: RNN

Recurrent Neural Network (RNN)


Input Sequence X 

Parameters A  B  C 


Activation Function ϕ : ReLU

Hidden States H_0 

Output Sequence Y 



Recurrent Neural Network (RNN)


Input Sequence X 

Parameters A  B  C 

Activation Function ϕ : ReLU



Hidden States H_0 

Output Sequence Y 




Exercise: LSTM

Initialize: Randomly set the previous hidden state h_0 to [1, 1] and memory cells C_0 to [0.3, -0.5]



Any question?



Readings for you:

- [Deep Learning book](#)
- [Forecasting \(FPP\) Book using Python](#)
- [AI by Hand by Tom Yeh](#)
- [Special thanks to Piyush Rai and Jay Alammar](#)
 - I adopted some of their slides available online.